

```
import pandas as pd
import seaborn as sb
import io
from google.colab import files
```

```
uploaded = files.upload()
df = pd.read_csv(io.BytesIO(uploaded["Auto.csv"]))
```

No file chosen

Upload widget is only available when the cell has been executed

in the current browser session. Please rerun this cell to enable.

Saving Auto.csv to Auto.csv

```
print(df.head)
print(df.shape)
```

```
<bound method NDFrame.head of      mpg  cylinders  displacement  horsepower  weight  acceleration  year  \
0    18.0         8      307.0         12.0  70.0
1    15.0         8      350.0         11.5  70.0
2    18.0         8      318.0         11.0  70.0
3    16.0         8      304.0         12.0  70.0
4    17.0         8      302.0         NaN  70.0
..  ...         ...         ...         ...         ...         ...
387  27.0         4      140.0         15.6  82.0
388  44.0         4       97.0         24.6  82.0
389  32.0         4      135.0         11.6  82.0
390  28.0         4      120.0         18.6  82.0
391  31.0         4      119.0         19.4  82.0
```

```
      origin      name
0         1  chevrolet chevelle malibu
1         1    buick skylark 320
2         1  plymouth satellite
3         1    amc rebel sst
4         1    ford torino
..  ...         ...
387        1    ford mustang gl
388         2      vw pickup
389         1  dodge rampage
390         1    ford ranger
391         1    chevy s-10
```

```
[392 rows x 9 columns]>
(392, 9)
```

```
print(df.dtypes)
```

```
mpg          float64
cylinders     int64
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year          float64
origin        int64
name          object
dtype: object
```

```
df.cylinders = df.cylinders.astype('category').cat.codes
df.origin = df.origin.astype('category')
print(df.dtypes)
```

```
mpg          float64
cylinders     int8
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year          float64
origin        category
name          object
dtype: object
```

```
df.dropna(inplace=True)
df.shape
```

```
(389, 9)
```

```
average = df['mpg'].mean()
df['mpg_high'] = df['mpg'] > average
df['mpg_high'] = df['mpg_high'].astype('int64')
df = df.drop('mpg', axis=1)
df = df.drop('name', axis=1)
df.head
```

```
<bound method NDFrame.head of
0      4      307.0      130      3504      12.0      70.0      1
1      4      350.0      165      3693      11.5      70.0      1
2      4      318.0      150      3436      11.0      70.0      1
3      4      304.0      150      3433      12.0      70.0      1
6      4      454.0      220      4354      9.0      70.0      1
..     ...      ...      ...      ...      ...      ...
387     1      140.0      86      2790      15.6      82.0      1
388     1       97.0      52      2130      24.6      82.0      2
389     1      135.0      84      2295      11.6      82.0      1
390     1      120.0      79      2625      18.6      82.0      1
391     1      119.0      82      2720      19.4      82.0      1
```

```
mpg_high
0      0
1      0
2      0
3      0
6      0
..     ...
387     1
388     1
389     1
390     1
391     1
```


```
[389 rows x 8 columns]>
```

```
print(df.index)
```

```
Int64Index([ 0, 1, 2, 3, 6, 7, 8, 9, 10, 11,
            ...,
            382, 383, 384, 385, 386, 387, 388, 389, 390, 391],
            dtype='int64', length=389)
```

```
sb.catplot(data=df, x=df.index, y='mpg_high')
```


```
<seaborn.axisgrid.FacetGrid at 0x7f837cd241c0>
```



Double-click (or enter) to edit


```
sb.relplot(data=df, x='horsepower', y='weight', hue='mpg_high')
```

```
<seaborn.axisgrid.FacetGrid at 0x7f837a18b610>
```



```
sb.boxplot(data=df, x='mpg_high', y='weight')
```

```
<Axes: xlabel='mpg_high', ylabel='weight'>
```



```
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
```

```
X_train, X_test, y_train, y_test = train_test_split(df.loc[:, df.columns != 'mpg_high'], df.mpg_high, test_size=0.20, random_state=1234)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(311, 7)
(78, 7)
(311,)
(78,)
```

```
lr_model = LogisticRegression(solver='lbfgs', max_iter=1000)
lr_model.fit(X_train, y_train)
print(lr_model.score(X_train, y_train))
```

```
0.9067524115755627
```

```
y_pred = lr_model.predict(X_test)
print(classification_report(y_pred, y_test))
```

	precision	recall	f1-score	support
0	0.80	0.98	0.88	41
1	0.96	0.73	0.83	37
accuracy			0.86	78
macro avg	0.88	0.85	0.85	78
weighted avg	0.88	0.86	0.86	78

```
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
```

```
dtree = DecisionTreeClassifier()
dtree.fit(X_train, y_train)
print(dtree.score(X_train, y_train))
```

1.0

```
y_pred = dtree.predict(X_test)
print(classification_report(y_pred, y_test))
```


	precision	recall	f1-score	support
0	0.88	0.98	0.93	45
1	0.96	0.82	0.89	33
accuracy			0.91	78
macro avg	0.92	0.90	0.91	78
weighted avg	0.92	0.91	0.91	78

```
tree.plot_tree(dtree)
```

```

[Text(0.6597222222222222, 0.9444444444444444, 'x[0] <= 2.5\ngini = 0.5\nsamples = 311\nvalue = [153, 158]'),
Text(0.4583333333333333, 0.8333333333333334, 'x[2] <= 101.0\ngini = 0.239\nsamples = 173\nvalue = [24, 149]'),
Text(0.3055555555555556, 0.7222222222222222, 'x[5] <= 75.5\ngini = 0.179\nsamples = 161\nvalue = [16, 145]'),
Text(0.1666666666666667, 0.6111111111111112, 'x[1] <= 119.5\ngini = 0.362\nsamples = 59\nvalue = [14, 45]'),
Text(0.0555555555555556, 0.5, 'x[4] <= 13.75\ngini = 0.159\nsamples = 46\nvalue = [4, 42]'),
Text(0.02777777777777778, 0.3888888888888889, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(0.08333333333333333, 0.3888888888888889, 'x[3] <= 2683.0\ngini = 0.087\nsamples = 44\nvalue = [2, 42]'),
Text(0.0555555555555556, 0.2777777777777778, 'x[3] <= 2377.0\ngini = 0.045\nsamples = 43\nvalue = [1, 42]'),
Text(0.02777777777777778, 0.1666666666666667, 'gini = 0.0\nsamples = 38\nvalue = [0, 38]'),
Text(0.08333333333333333, 0.1666666666666667, 'x[3] <= 2385.0\ngini = 0.32\nsamples = 5\nvalue = [1, 4]'),
Text(0.0555555555555556, 0.0555555555555556, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(0.1111111111111111, 0.0555555555555556, 'gini = 0.0\nsamples = 4\nvalue = [0, 4]'),
Text(0.1111111111111111, 0.2777777777777778, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(0.2777777777777778, 0.5, 'x[3] <= 2567.0\ngini = 0.355\nsamples = 13\nvalue = [10, 3]'),
Text(0.25, 0.3888888888888889, 'x[3] <= 2429.5\ngini = 0.469\nsamples = 8\nvalue = [5, 3]'),
Text(0.2222222222222222, 0.2777777777777778, 'x[4] <= 15.75\ngini = 0.278\nsamples = 6\nvalue = [5, 1]'),
Text(0.1944444444444444, 0.1666666666666667, 'x[4] <= 14.75\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(0.1666666666666667, 0.0555555555555556, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(0.2222222222222222, 0.0555555555555556, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(0.25, 0.1666666666666667, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
Text(0.2777777777777778, 0.2777777777777778, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(0.3055555555555556, 0.3888888888888889, 'gini = 0.0\nsamples = 5\nvalue = [5, 0]'),
Text(0.4444444444444444, 0.6111111111111112, 'x[3] <= 3250.0\ngini = 0.038\nsamples = 102\nvalue = [2, 100]'),
Text(0.3888888888888889, 0.5, 'x[3] <= 2880.0\ngini = 0.02\nsamples = 100\nvalue = [1, 99]'),
Text(0.3611111111111111, 0.3888888888888889, 'gini = 0.0\nsamples = 94\nvalue = [0, 94]'),
Text(0.4166666666666667, 0.3888888888888889, 'x[3] <= 2920.0\ngini = 0.278\nsamples = 6\nvalue = [1, 5]'),
Text(0.3888888888888889, 0.2777777777777778, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(0.4444444444444444, 0.2777777777777778, 'gini = 0.0\nsamples = 5\nvalue = [0, 5]'),
Text(0.5, 0.5, 'x[5] <= 77.5\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(0.4722222222222222, 0.3888888888888889, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(0.5277777777777778, 0.3888888888888889, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(0.6111111111111112, 0.7222222222222222, 'x[4] <= 14.45\ngini = 0.444\nsamples = 12\nvalue = [8, 4]'),
Text(0.5833333333333334, 0.6111111111111112, 'x[5] <= 76.0\ngini = 0.444\nsamples = 6\nvalue = [2, 4]'),
Text(0.5555555555555556, 0.5, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),
Text(0.6111111111111112, 0.5, 'x[1] <= 138.5\ngini = 0.444\nsamples = 3\nvalue = [2, 1]'),
Text(0.5833333333333334, 0.3888888888888889, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(0.6388888888888889, 0.3888888888888889, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(0.6388888888888889, 0.6111111111111112, 'gini = 0.0\nsamples = 6\nvalue = [6, 0]'),
Text(0.8611111111111112, 0.8333333333333334, 'x[5] <= 79.5\ngini = 0.122\nsamples = 138\nvalue = [129, 9]'),
Text(0.8055555555555556, 0.7222222222222222, 'x[4] <= 21.6\ngini = 0.045\nsamples = 129\nvalue = [126, 3]'),
Text(0.7777777777777778, 0.6111111111111112, 'x[3] <= 2737.0\ngini = 0.031\nsamples = 128\nvalue = [126, 2]'),
Text(0.7222222222222222, 0.5, 'x[5] <= 76.0\ngini = 0.444\nsamples = 3\nvalue = [2, 1]'),
Text(0.6944444444444444, 0.3888888888888889, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(0.75, 0.3888888888888889, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(0.8333333333333334, 0.5, 'x[2] <= 83.0\ngini = 0.016\nsamples = 125\nvalue = [124, 1]'),
Text(0.8055555555555556, 0.3888888888888889, 'x[1] <= 225.0\ngini = 0.375\nsamples = 4\nvalue = [3, 1]'),
Text(0.7777777777777778, 0.2777777777777778, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(0.8333333333333334, 0.2777777777777778, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),
Text(0.8611111111111112, 0.3888888888888889, 'gini = 0.0\nsamples = 121\nvalue = [121, 0]'),
Text(0.8333333333333334, 0.6111111111111112, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(0.9166666666666667, 0.7222222222222222, 'x[1] <= 196.5\ngini = 0.444\nsamples = 9\nvalue = [3, 6]'),
Text(0.8888888888888889, 0.6111111111111112, 'gini = 0.0\nsamples = 4\nvalue = [0, 4]'),
Text(0.9444444444444444, 0.6111111111111112, 'x[1] <= 247.0\ngini = 0.48\nsamples = 5\nvalue = [3, 2]'),
Text(0.9166666666666667, 0.5, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),
Text(0.9722222222222222, 0.5, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]')]

```



```
from sklearn.neural_network import MLPClassifier
```

```
model_1 = MLPClassifier(hidden_layer_sizes=(100, 50), activation='logistic')
model_1.fit(X_train, y_train)
y_pred = model_1.predict(X_test)
print(classification_report(y_pred, y_test))
```

	precision	recall	f1-score	support
0	0.84	0.98	0.90	43
1	0.96	0.77	0.86	35
accuracy			0.88	78
macro avg	0.90	0.87	0.88	78
weighted avg	0.90	0.88	0.88	78

```
/usr/local/lib/python3.9/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimizer:
warnings.warn(
```



```
model_2 = MLPClassifier() #default network with hidden layer size (100,) and relu activatio function
model_2.fit(X_train, y_train)
y_pred = model_2.predict(X_test)
print(classification_report(y_pred, y_test))
```

	precision	recall	f1-score	support
0	0.84	0.98	0.90	43
1	0.96	0.77	0.86	35
accuracy			0.88	78
macro avg	0.90	0.87	0.88	78
weighted avg	0.90	0.88	0.88	78

Neural Net Outcomes

The first neural net with 2 hidden layers (100, 50) and a logistic activation function worked better than the default MLPClassifier

Overall Analysis

The most successful model in terms of accuracy is the Decision Tree Classifier. The Logistic Regression Classifier and first Neural Networks performed about the same with logistic regression. The default neural net performed the worst. While the recall of precision of each model is relatively equal, the accuracy has quite a large spread. from 79% to 91%

This data set is pretty small, so it makes sense that the neural nets did not perform better. Since the logistic regression model performed worse, it may be the case the is not simply linearly separable and a non linear clasifier (Decision Tree) is necessary.

I prefer python to R because it is more programmer friendly. I also think the syntax is easier to comprehend. Even though R tries to be verbose to benefit non programmers, I think this is a situation where less is more.

