

```

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
import matplotlib.pyplot as plt
# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all files under the input directory

# import os
# for dirname, _, filenames in os.walk('/kaggle/input'):
#     for filename in filenames:
#         print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets preserved as output when you create a version using "Save &
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside of the current session

image_size = (256, 256)
batch_size = 1
train_ds = tf.keras.utils.image_dataset_from_directory(
    '/kaggle/input/garbage-classification/Garbage classification/Garbage classification',
    validation_split=0.2,
    subset="training",
    image_size=image_size,
    seed=123,
    batch_size=batch_size,
    color_mode="rgb"
)
val_ds = tf.keras.utils.image_dataset_from_directory(
    '/kaggle/input/garbage-classification/Garbage classification/Garbage classification',
    validation_split=0.2,
    subset="validation",
    image_size=image_size,
    seed=123,
    batch_size=batch_size,
    color_mode="rgb"
)

 Found 2527 files belonging to 6 classes.
Using 2022 files for training.
Found 2527 files belonging to 6 classes.
Using 505 files for validation.

class_names = train_ds.class_names
num_classes = len(train_ds.class_names)
image_shape = image_size + (3,)
count = dict.fromkeys(class_names, 0)
for image, labels in train_ds:
    for i in range(batch_size):
        count[class_names[labels[0]]] += 1

plt.bar(range(len(count)), list(count.values()), align='center')
plt.xticks(range(len(count)), list(count.keys()))

```

```
([<matplotlib.axis.XTick at 0x7d8d70411750>,
<matplotlib.axis.XTick at 0x7d8d704113d0>,
<matplotlib.axis.XTick at 0x7d8d703f1f10>,
<matplotlib.axis.XTick at 0x7d8d70381690>,
<matplotlib.axis.XTick at 0x7d8d70381710>,
<matplotlib.axis.XTick at 0x7d8d703813d0>],
[Text(0, 0, 'cardboard'),
Text(1, 0, 'glass'),
Text(2, 0, 'metal'),
Text(3, 0, 'paper'),
Text(4, 0, 'plastic'),
Text(5, 0, 'trash')])
```




```
plt.figure(figsize=(10,10))
```

```
for i, example in enumerate(train_ds.take(9)):
    image = example[0][0]
    label = example[1][0]
```

```
ax = plt.subplot(3, 3, i+1)
plt.imshow(image.numpy().astype("uint8"))
plt.title(class_names[label])
plt.axis("off")
```

metal



glass



paper



glass



glass



paper



metal



cardboard



glass



Neural Network

```

model = Sequential([
    layers.Input(shape=image_shape),
    layers.Dense(128, activation='relu'),
    layers.Dropout(.2),
    layers.Dense(256, activation='relu'),
    layers.Dropout(.2),
    layers.Flatten(),
    layers.Dense(num_classes)
])

```

```

model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

```

```

epochs=15
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs
)

```

```

Epoch 1/15
2023-04-22 22:03:20.388929: E tensorflow/core/grappler/optimizers/meta_optimizer.cc:954] layout failed: INVALID_ARGUMENT: Size of values
2022/2022 [=====] - 85s 39ms/step - loss: 542.3392 - accuracy: 0.2834 - val_loss: 1.8292 - val_accuracy: 0.0574
Epoch 2/15
2022/2022 [=====] - 78s 39ms/step - loss: 1.7465 - accuracy: 0.3294 - val_loss: 1.8465 - val_accuracy: 0.1723
Epoch 3/15
2022/2022 [=====] - 78s 39ms/step - loss: 1.8629 - accuracy: 0.2982 - val_loss: 1.8199 - val_accuracy: 0.1743
Epoch 4/15
2022/2022 [=====] - 78s 39ms/step - loss: 1.7995 - accuracy: 0.2794 - val_loss: 1.8029 - val_accuracy: 0.1762
Epoch 5/15
2022/2022 [=====] - 78s 39ms/step - loss: 1.7466 - accuracy: 0.2621 - val_loss: 1.7660 - val_accuracy: 0.1822
Epoch 6/15
2022/2022 [=====] - 78s 39ms/step - loss: 1.8256 - accuracy: 0.2463 - val_loss: 1.7435 - val_accuracy: 0.2475
Epoch 7/15
2022/2022 [=====] - 78s 39ms/step - loss: 1.7273 - accuracy: 0.2730 - val_loss: 1.7201 - val_accuracy: 0.2653
Epoch 8/15
2022/2022 [=====] - 78s 39ms/step - loss: 1.6699 - accuracy: 0.3205 - val_loss: 1.9330 - val_accuracy: 0.2812
Epoch 9/15
2022/2022 [=====] - 78s 39ms/step - loss: 1.6018 - accuracy: 0.3398 - val_loss: 2.4155 - val_accuracy: 0.2752
Epoch 10/15
2022/2022 [=====] - 78s 39ms/step - loss: 1.5030 - accuracy: 0.3961 - val_loss: 3.7635 - val_accuracy: 0.2792
Epoch 11/15
2022/2022 [=====] - 78s 39ms/step - loss: 1.5932 - accuracy: 0.4085 - val_loss: 3.4495 - val_accuracy: 0.2792
Epoch 12/15
2022/2022 [=====] - 78s 39ms/step - loss: 1.5424 - accuracy: 0.4095 - val_loss: 2.2216 - val_accuracy: 0.2772
Epoch 13/15
2022/2022 [=====] - 78s 39ms/step - loss: 1.4318 - accuracy: 0.4426 - val_loss: 2.7552 - val_accuracy: 0.2970
Epoch 14/15
2022/2022 [=====] - 78s 39ms/step - loss: 1.2871 - accuracy: 0.4679 - val_loss: 4.6169 - val_accuracy: 0.3188
Epoch 15/15
2022/2022 [=====] - 78s 39ms/step - loss: 1.3463 - accuracy: 0.4703 - val_loss: 4.3838 - val_accuracy: 0.3188

```



```

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

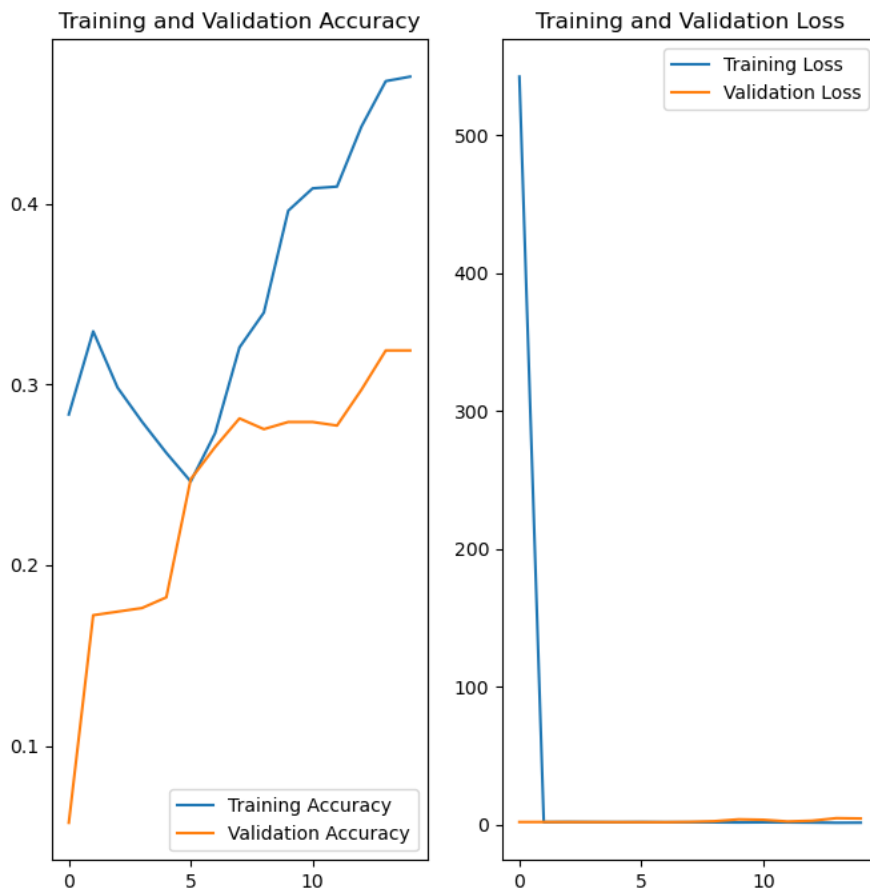
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(epochs)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()

```



▼ Convolutional Neural Network

```

model = Sequential([
    layers.Rescaling(1./255, input_shape=image_shape),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])

```

```
model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 256, 256, 3)	0
conv2d (Conv2D)	(None, 256, 256, 16)	448
max_pooling2d (MaxPooling2D)	(None, 128, 128, 16)	0
conv2d_1 (Conv2D)	(None, 128, 128, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 64, 64, 32)	0
conv2d_2 (Conv2D)	(None, 64, 64, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 32, 32, 64)	0
flatten_1 (Flatten)	(None, 65536)	0
dense_3 (Dense)	(None, 128)	8388736
dense_4 (Dense)	(None, 6)	774

Total params: 8,413,094
Trainable params: 8,413,094
Non-trainable params: 0

```
epochs=15
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs
)

Epoch 1/15
2022/2022 [=====] - 25s 8ms/step - loss: 1.7200 - accuracy: 0.2591 - val_loss: 1.5816 - val_accuracy: 0.3366
Epoch 2/15
2022/2022 [=====] - 14s 7ms/step - loss: 1.5269 - accuracy: 0.3813 - val_loss: 1.5331 - val_accuracy: 0.3564
Epoch 3/15
2022/2022 [=====] - 15s 7ms/step - loss: 1.2378 - accuracy: 0.5178 - val_loss: 1.5599 - val_accuracy: 0.3584
Epoch 4/15
2022/2022 [=====] - 16s 8ms/step - loss: 0.8047 - accuracy: 0.7018 - val_loss: 1.9094 - val_accuracy: 0.4198
Epoch 5/15
2022/2022 [=====] - 15s 7ms/step - loss: 0.4572 - accuracy: 0.8432 - val_loss: 2.5062 - val_accuracy: 0.3881
Epoch 6/15
2022/2022 [=====] - 15s 7ms/step - loss: 0.2998 - accuracy: 0.9105 - val_loss: 3.5812 - val_accuracy: 0.4238
Epoch 7/15
2022/2022 [=====] - 15s 7ms/step - loss: 0.1926 - accuracy: 0.9357 - val_loss: 4.2626 - val_accuracy: 0.4059
Epoch 8/15
2022/2022 [=====] - 15s 7ms/step - loss: 0.1566 - accuracy: 0.9575 - val_loss: 5.0650 - val_accuracy: 0.4198
Epoch 9/15
2022/2022 [=====] - 15s 7ms/step - loss: 0.1156 - accuracy: 0.9674 - val_loss: 5.8556 - val_accuracy: 0.3683
Epoch 10/15
2022/2022 [=====] - 15s 7ms/step - loss: 0.1141 - accuracy: 0.9683 - val_loss: 5.5162 - val_accuracy: 0.4000
Epoch 11/15
2022/2022 [=====] - 15s 7ms/step - loss: 0.0589 - accuracy: 0.9871 - val_loss: 6.2815 - val_accuracy: 0.4020
Epoch 12/15
2022/2022 [=====] - 14s 7ms/step - loss: 0.0412 - accuracy: 0.9901 - val_loss: 8.1907 - val_accuracy: 0.3743
Epoch 13/15
2022/2022 [=====] - 14s 7ms/step - loss: 0.0563 - accuracy: 0.9866 - val_loss: 7.3596 - val_accuracy: 0.4198
Epoch 14/15
2022/2022 [=====] - 15s 7ms/step - loss: 0.0627 - accuracy: 0.9847 - val_loss: 7.6886 - val_accuracy: 0.3703
Epoch 15/15
2022/2022 [=====] - 15s 7ms/step - loss: 0.0836 - accuracy: 0.9812 - val_loss: 7.3248 - val_accuracy: 0.4099

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
```

```


loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(epochs)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()

```



▼ Transfer Learning

```
preprocess_input = tf.keras.applications.mobilenet_v2.preprocess_input
```

▼ Base Model

```

# Create the base model from the pre-trained model MobileNet V2
base_model = tf.keras.applications.MobileNetV2(input_shape=image_shape,
                                                include_top=False,
                                                weights='imagenet')

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_9406464/9406464 [=====] - 0s 0us/step

```
image_batch, label_batch = next(iter(train_ds))
feature_batch = base_model(image_batch)
print(feature_batch.shape)
```

```
(1, 8, 8, 1280)
```

```
base_model.trainable = False
```

▼ Feature Extraction

▼ Data Augmentation

```
data_augmentation = tf.keras.Sequential([
    tf.keras.layers.RandomFlip('horizontal'),
    tf.keras.layers.RandomRotation(0.2),
])
```

▼ Adding Classification Head

```
global_average_layer = tf.keras.layers.GlobalAveragePooling2D()
feature_batch_average = global_average_layer(feature_batch)
print(feature_batch_average.shape)
```

```
(1, 1280)
```

```
prediction_layer = tf.keras.layers.Dense(1)
prediction_batch = prediction_layer(feature_batch_average)
print(prediction_batch.shape)
```

```
(1, 1)
```

```
inputs = tf.keras.Input(shape=image_shape)
x = data_augmentation(inputs)
x = preprocess_input(x)
x = base_model(x, training=False)
x = global_average_layer(x)
x = tf.keras.layers.Dropout(0.2)(x)
outputs = prediction_layer(x)
model = tf.keras.Model(inputs, outputs)
```

```
base_learning_rate = 0.0001
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=base_learning_rate),
              loss=tf.keras.losses.BinaryCrossentropy(from_logits=True),
              metrics=['accuracy'])
```

```
model.summary()
```

```
Model: "model"
```

Layer (type)	Output Shape	Param #
input_3 (InputLayer)	[(None, 256, 256, 3)]	0
sequential_2 (Sequential)	(None, 256, 256, 3)	0
tf.math.truediv (TFOpLambda)	(None, 256, 256, 3)	0
tf.math.subtract (TFOpLambda)	(None, 256, 256, 3)	0
mobilenetv2_1.00_224 (Functional)	(None, 8, 8, 1280)	2257984
global_average_pooling2d (G	(None, 1280)	0

```

    globalAveragePooling2D)

dropout_2 (Dropout)          (None, 1280)          0

dense_5 (Dense)              (None, 1)             1281

=====
Total params: 2,259,265
Trainable params: 1,281
Non-trainable params: 2,257,984

```

```

initial_epochs = 10

loss0, accuracy0 = model.evaluate(val_ds)

505/505 [=====] - 6s 9ms/step - loss: -0.2249 - accuracy: 0.2317


print("initial loss: {:.2f}".format(loss0))
print("initial accuracy: {:.2f}".format(accuracy0))

initial loss: -0.22
initial accuracy: 0.23

history = model.fit(train_ds,
                    epochs=initial_epochs,
                    validation_data=val_ds)

Epoch 1/10
2022/2022 [=====] - 35s 15ms/step - loss: -21.1117 - accuracy: 0.1973 - val_loss: -39.1773 - val_accuracy: 0.20
Epoch 2/10
2022/2022 [=====] - 30s 15ms/step - loss: -61.1551 - accuracy: 0.1968 - val_loss: -77.5614 - val_accuracy: 0.20
Epoch 3/10
2022/2022 [=====] - 30s 15ms/step - loss: -102.3228 - accuracy: 0.1968 - val_loss: -116.0098 - val_accuracy: 0.
Epoch 4/10
2022/2022 [=====] - 30s 15ms/step - loss: -142.5291 - accuracy: 0.1968 - val_loss: -154.5656 - val_accuracy: 0.
Epoch 5/10
2022/2022 [=====] - 30s 15ms/step - loss: -183.1707 - accuracy: 0.1968 - val_loss: -193.0489 - val_accuracy: 0.
Epoch 6/10
2022/2022 [=====] - 30s 15ms/step - loss: -223.3655 - accuracy: 0.1968 - val_loss: -231.4556 - val_accuracy: 0.
Epoch 7/10
2022/2022 [=====] - 30s 15ms/step - loss: -264.3979 - accuracy: 0.1968 - val_loss: -269.9103 - val_accuracy: 0.
Epoch 8/10
2022/2022 [=====] - 30s 15ms/step - loss: -304.4212 - accuracy: 0.1968 - val_loss: -308.3747 - val_accuracy: 0.
Epoch 9/10
2022/2022 [=====] - 30s 15ms/step - loss: -344.6414 - accuracy: 0.1968 - val_loss: -346.7709 - val_accuracy: 0.
Epoch 10/10
2022/2022 [=====] - 30s 15ms/step - loss: -384.1013 - accuracy: 0.1968 - val_loss: -385.1770 - val_accuracy: 0.

```



```

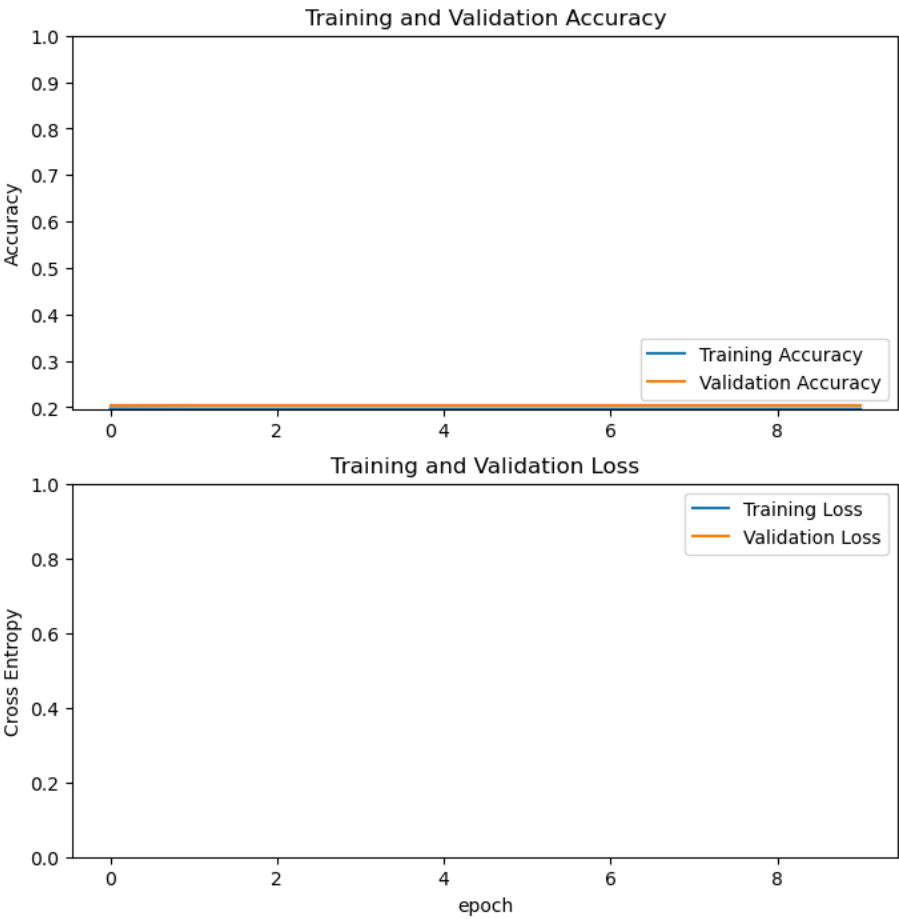
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

loss = history.history['loss']
val_loss = history.history['val_loss']

plt.figure(figsize=(8, 8))
plt.subplot(2, 1, 1)
plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.ylabel('Accuracy')
plt.ylim([min(plt.ylim()), max(plt.ylim())])
plt.title('Training and Validation Accuracy')

plt.subplot(2, 1, 2)
plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.ylabel('Cross Entropy')
plt.ylim([min(plt.ylim()), 1.0])
plt.title('Training and Validation Loss')
plt.xlabel('epoch')
plt.show()

```

▼ Fine Tuning

```
base_model.trainable = True

# Let's take a look to see how many layers are in the base model
print("Number of layers in the base model: ", len(base_model.layers))

# Fine-tune from this layer onwards
fine_tune_at = 100

# Freeze all the layers before the `fine_tune_at` layer
for layer in base_model.layers[:fine_tune_at]:
    layer.trainable = False

Number of layers in the base model: 154

model.compile(loss=tf.keras.losses.BinaryCrossentropy(from_logits=True),
              optimizer = tf.keras.optimizers.RMSprop(learning_rate=base_learning_rate/10),
              metrics=['accuracy'])
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_3 (InputLayer)	[(None, 256, 256, 3)]	0
sequential_2 (Sequential)	(None, 256, 256, 3)	0
tf.math.truediv (TFOpLambda)	(None, 256, 256, 3)	0
tf.math.subtract (TFOpLambda)	(None, 256, 256, 3)	0
mobilenetv2_1.00_224 (Functional)	(None, 8, 8, 1280)	2257984

```

global_average_pooling2d (G   (None, 1280)      0
lobalAveragePooling2D)

dropout_2 (Dropout)          (None, 1280)      0

dense_5 (Dense)              (None, 1)         1281

```

```

=====
Total params: 2,259,265
Trainable params: 1,862,721
Non-trainable params: 396,544

```

```

fine_tune_epochs = 10
total_epochs = initial_epochs + fine_tune_epochs

```

```

history_fine = model.fit(train_ds,
                          epochs=total_epochs,
                          initial_epoch=history.epoch[-1],
                          validation_data=val_ds)

```

```

Epoch 10/20
2022/2022 [=====] - 55s 22ms/step - loss: -6343.2456 - accuracy: 0.1968 - val_loss: -6687.9141 - val_accuracy:
Epoch 11/20
2022/2022 [=====] - 44s 22ms/step - loss: -6887.1162 - accuracy: 0.1968 - val_loss: -6748.1274 - val_accuracy:
Epoch 12/20
2022/2022 [=====] - 44s 22ms/step - loss: -7047.5200 - accuracy: 0.1968 - val_loss: -7014.0615 - val_accuracy:
Epoch 13/20
2022/2022 [=====] - 44s 22ms/step - loss: -7208.3350 - accuracy: 0.1968 - val_loss: -7154.0527 - val_accuracy:
Epoch 14/20
2022/2022 [=====] - 44s 22ms/step - loss: -7354.4229 - accuracy: 0.1968 - val_loss: -7254.2358 - val_accuracy:
Epoch 15/20
2022/2022 [=====] - 44s 22ms/step - loss: -7507.6934 - accuracy: 0.1968 - val_loss: -7436.2754 - val_accuracy:
Epoch 16/20
2022/2022 [=====] - 44s 22ms/step - loss: -7667.2451 - accuracy: 0.1968 - val_loss: -7540.5107 - val_accuracy:
Epoch 17/20
2022/2022 [=====] - 44s 22ms/step - loss: -7819.6914 - accuracy: 0.1968 - val_loss: -7678.7139 - val_accuracy:
Epoch 18/20
2022/2022 [=====] - 44s 22ms/step - loss: -7972.3208 - accuracy: 0.1968 - val_loss: -7847.5215 - val_accuracy:
Epoch 19/20
2022/2022 [=====] - 44s 22ms/step - loss: -8113.9160 - accuracy: 0.1968 - val_loss: -7954.1084 - val_accuracy:
Epoch 20/20
2022/2022 [=====] - 44s 22ms/step - loss: -8250.7617 - accuracy: 0.1968 - val_loss: -8135.7822 - val_accuracy:

```

```

acc += history_fine.history['accuracy']
val_acc += history_fine.history['val_accuracy']

```

```

loss += history_fine.history['loss']
val_loss += history_fine.history['val_loss']

```

```

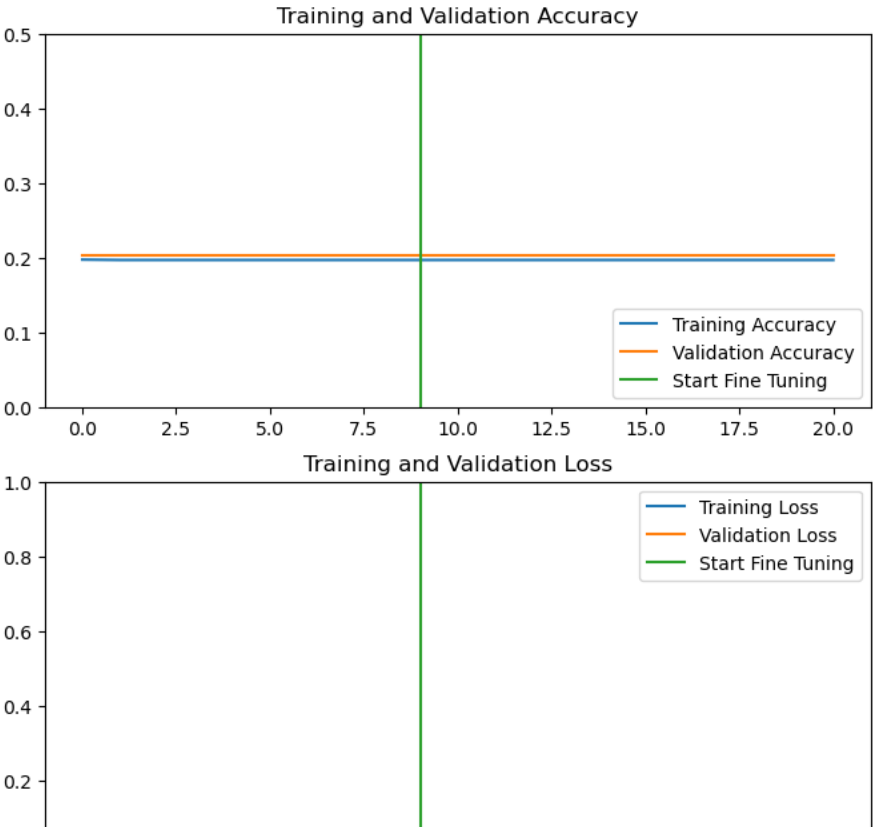
plt.figure(figsize=(8, 8))
plt.subplot(2, 1, 1)
plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.ylim([0, .5])
plt.plot([initial_epochs-1,initial_epochs-1],
         plt.ylim(), label='Start Fine Tuning')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

```

```

plt.subplot(2, 1, 2)
plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.ylim([0, 1.0])
plt.plot([initial_epochs-1,initial_epochs-1],
         plt.ylim(), label='Start Fine Tuning')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.xlabel('epoch')
plt.show()

```



Analysis